


(REMINDER)

DAY 067

What is RAG?



(REMINDER)



In one line

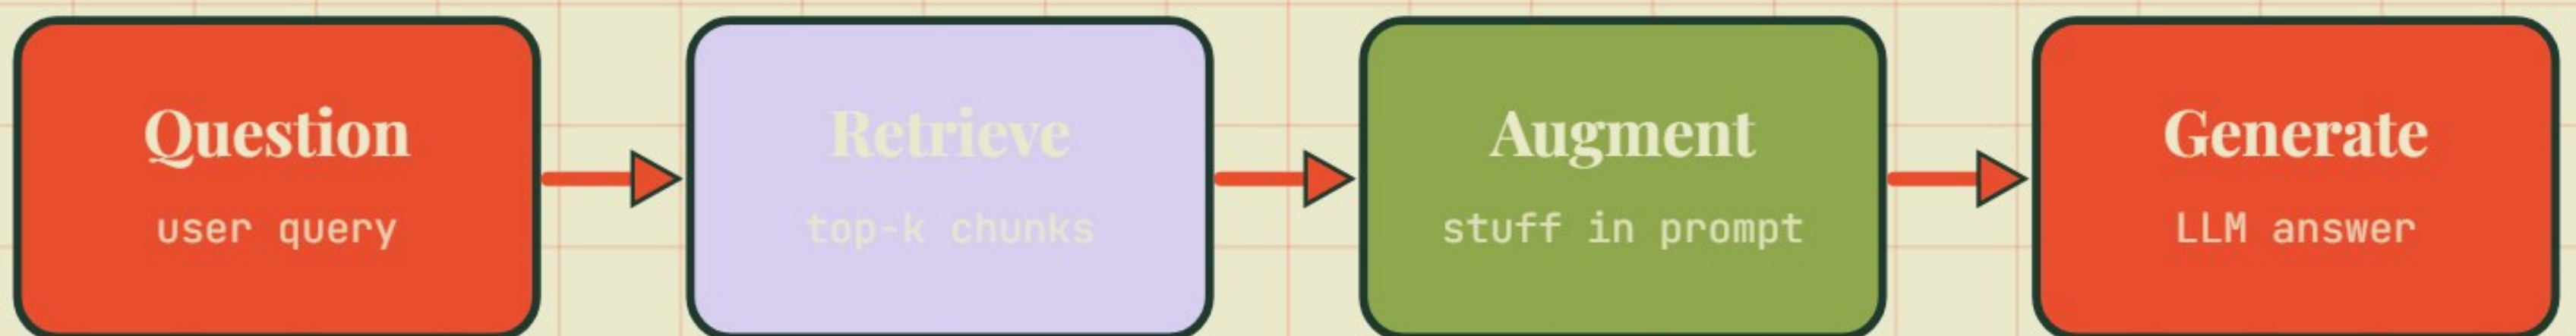
Retrieval-Augmented Generation lets an LLM look up info from your documents before answering — so it stops making things up.



(REMINDER)



The RAG flow



Why bother?

- 1 Fresh data without retraining
- 2 Real citations
- 3 Cheaper than fine-tuning
- 4 Works on private docs

(REMINDER)



Mental model

Imagine an open-book exam. The LLM is the student.
RAG is what hands it the right pages — right when it
needs them.



Chatbot vs RAG

Plain LLM

- Frozen knowledge
- No citations
- Hallucinates
- Hard to update

RAG

- Live data
- Real citations
- Grounded
- Easy to update

(REMINDER)



Stack to try

- 1 LlamaIndex or LangChain
- 2 OpenAI / Cohere embeddings
- 3 Chroma / Qdrant / pgvector
- 4 Cohere or BGE reranker



(REMINDER)



DAY 067

Save this.
Follow for
more.

